

# An Algebra of Prediction-Rewarding Mechanisms

Scoring rules, market makers, and pools as composable transducers

Peter Cotton

Working draft v0.3 · 2026

## Abstract

Scoring rules, market makers, parimutuels, and opinion pools compose when each stage is a stateful transducer over one message type, and the message deployed contests actually collect is a finite cloud of samples. This note provides sample-based elicitation inside that algebra: jittered settlement makes cloud submission strictly proper, the composition operators act on clouds pointwise, and joint laws factor into margin stages plus a copula stage settled on the rank vector. The single-stage theory underneath is classical convex duality: proper scores from convex entropies, cost-function market makers by Fenchel conjugacy, the two opinion pools as the two Kullback-Leibler barycenters, merged market makers as the infimal convolution of risk sharing. The benefit of composing is illustrated by the simplest two-stage example, where stagewise equilibrium dominates conformal prediction: a residual market collects the conformal predictor’s information gap  $I(R; X)$  as bankroll growth while its marginal coverage stays exact.

## 1 Introduction

Mechanisms for eliciting and aggregating forecasts are usually studied one at a time. A proper scoring rule is analysed as a one-shot contract; a market maker as a sequential trading venue; an opinion pool as an estimator; a calibration test as a diagnostic. Deployed systems are thinner than the theory allows: with few exceptions they run one level of competition against one internal model. Numerai pools staked submissions into a single stake-weighted meta-model and pays each forecast its marginal contribution to it ([Craib et al. 2017](#)); CrunchDAO blends each contest into one ensemble; the IARPA prediction polls aggregated forecasts by track record, comparing favourably with head-to-head markets ([Atanasov et al. 2017](#)), with reputation rather than wealth as the threaded state. In every case, one pool and one internal aggregate.

Chained mechanisms are old and live in adjacent markets, priced in points rather than distributions. Two-settlement electricity markets settle real-time deviations from day-ahead positions, so the second stage prices the residual of the first, and convergence bidding pays financial forecasters for correcting the day-ahead consensus ([Cramton 2017](#); [Jha and Wolak 2023](#)). Racetrack totes have priced margins and joint finishing orders in separate simultaneous pools since 1931: the win pool carries the marginals, the exotic pools the dependence, and the consistency between them supported documented arbitrage ([Harville 1973](#); [Hausch et al. 1981](#)). What is scarce is the composed design for *distributional* forecasts. The only live, long-running platform we are aware of that composes past one stage is monteprediction, a weekly self-funding pool over million-scenario joint submissions in eleven dimensions, with wealth threaded across rounds since January 2024 ([Cotton 2024b](#)). The full chain ran on the microprediction platform

(Cotton 2022): streams spawned z-streams of community percentiles, and bivariate and trivariate dependence streams priced copulas, so calibration and dependence were themselves the subject of further games. The stacked-lottery design behind it was presented at MIT CSAIL in 2020 (Cotton 2020, slides 29-31); that platform is retired.

Chaining stages raises the question of what each stage elicits, and the answer moves under transformation of the message or the outcome. Scoring a kernel-smoothed submission at the raw outcome elicits the deconvolution of the belief rather than the belief; jittering the settlement by the smoothing kernel repairs it (Cotton 2026a). This is a property of a single stage, not of the chain, and it is the third question below. The repair makes a finite sample cloud a legitimate message type, and §6 runs the algebra on the objects deployed contests actually collect.

Conformal prediction supplies a second motivation, from the opposite direction. Split-conformal is itself a composition, a point predictor chained into a rank-based calibration stage, but a degenerate one: the pool step is skipped by fiat, all credit assigned to a single model in advance, and the calibration stage prices the residual flat in the input. The wasted log score of a single-shape conformal predictor is the mutual information  $I(R; X)$  between residual and input, and an entrant to a parimutuel residual pool who conditions on the input collects it as bankroll growth at exactly that rate (Theorem 9). Running the residual stage as an actual pool, rather than assuming its winner, is what the operators below are for; Proposition 4 locates the gap in the theory of the probability integral transform, and §5 prices it.

This note organizes the catalogue around three questions that are often blurred together:

1. *Message closure.* Can every stage consume and emit one common object, so that stages plug together at all?
2. *Convex generation.* Is each individual mechanism generated by a convex potential, so that one dictionary covers scoring rules, market makers, and their duals?
3. *Propriety under transformation.* When a stage transforms forecasts or outcomes, does strict propriety survive?

The first is an engineering convention (§2, §8). The second and third admit theorems, which are stated with complete proofs in §3 and §4; that mathematics is classical, collected in one place and one convention. The contributions are the parts built on it: sample-based elicitation, running the algebra on clouds rather than exact densities (§6); the factoring of multivariate elicitation into margin stages and a rank-settled copula stage (§7); and the priced conformal example (§5), where the residual market's bankroll growth is the conformal predictor's information gap. Play is stagewise throughout: no participant deviates across stages (§4). Cross-stage strategy is outside the note's scope.

## 2 Preliminaries

**The finite setting.** Sections 3 and 4 work with a finite outcome set  $\{1, \dots, n\}$  and reports  $p$  in the probability simplex  $\Delta$ . The continuous theory uses proper scoring rules on probability measures (log score, CRPS, energy and kernel scores) or discretized approximations, and the finite statements below transfer with the usual measure-theoretic care; the one place where the transfer is not routine is the probability integral transform, treated separately in Proposition 4.

**Scores.** A scoring rule assigns  $S(p, i) \in \mathbb{R} \cup \{-\infty\}$  to a report  $p$  and outcome  $i$ ; the expected score of reporting  $p$  under belief  $q$  is  $S(p; q) = \sum_i q_i S(p, i)$ , assumed well defined with  $S(q; q)$

finite (this is the regularity used throughout). The rule is *proper* if  $S(q; q) \geq S(p; q)$  for all  $p, q$  and *strictly proper* if equality forces  $p = q$ . Scores are written in reward form: higher is better.

**Convex tools.** For convex  $G$  with subgradient selection  $G'(p) \in \partial G(p)$ , the Bregman divergence is  $D_G(q, p) = G(q) - G(p) - \langle G'(p), q - p \rangle$ . The Legendre-Fenchel conjugate of a function  $R$  on  $\Delta$  is  $R^*(\mathbf{q}) = \sup_{p \in \Delta} \langle p, \mathbf{q} \rangle - R(p)$ . The infimal convolution of  $f$  and  $g$  is  $(f \square g)(x) = \inf_y f(y) + g(x - y)$ . “Closed proper convex” is used in the standard sense (Rockafellar 1970). Differentiable statements are on the relative interior of  $\Delta$ , with extended-real values allowed on the boundary.

### 3 One potential, three mechanisms

**Theorem 1 (characterisation of proper scoring rules; Savage (1971), McCarthy (1956)).** *A regular scoring rule  $S$  is proper iff there is a convex  $G : \Delta \rightarrow \mathbb{R}$  with*

$$S(p, i) = G(p) + \langle G'(p), e_i - p \rangle, \quad G'(p) \in \partial G(p),$$

where  $e_i$  is the  $i$ -th unit vector. It is strictly proper iff  $G$  is strictly convex relative to  $\Delta$ , and then  $G(p) = S(p; p)$  is the expected score of a truthful forecaster.

**Proof.** ( $\Leftarrow$ ) Since  $\sum_i q_i(e_i - p) = q - p$ , the expected score is affine in the belief:  $S(p; q) = G(p) + \langle G'(p), q - p \rangle$ . Hence

$$S(q; q) - S(p; q) = G(q) - G(p) - \langle G'(p), q - p \rangle = D_G(q, p),$$

which is non-negative by the supporting-hyperplane inequality, and zero only at  $q = p$  when  $G$  is strictly convex relative to  $\Delta$ ; so  $S$  is (strictly) proper.

( $\Rightarrow$ ) Define  $G(q) := S(q; q)$ . Properness says  $G(q) = \sup_p S(p; q)$ , and each  $q \mapsto S(p; q) = \langle S(p, \cdot), q \rangle$  is affine, so  $G$  is a pointwise supremum of affine functions, hence convex. For any fixed  $p$ ,

$$S(p; q) = G(p) + \langle S(p, \cdot), q - p \rangle \quad \text{with} \quad S(p; q) \leq G(q) \quad \forall q,$$

so the vector  $S(p, \cdot)$  is a subgradient of  $G$  at  $p$  (acting on the tangent space of  $\Delta$ , where subgradients are defined up to a constant shift along  $\mathbf{1}$ , which the representation absorbs). Substituting  $G'(p) = S(p, \cdot)$  into the display returns  $S(p, i)$  identically. If  $S$  is strictly proper the supremum has a unique maximiser at every  $q$ , which for a supremum of affine functions is equivalent to strict convexity relative to  $\Delta$ . ■

The content of Theorem 1 is the dictionary *proper scoring rule*  $\leftrightarrow$  *convex function*  $\leftrightarrow$  *Bregman divergence* (Gneiting and Raftery 2007; Banerjee et al. 2005). The classics are three choices of  $G$ , with the divergences computed directly from the definition:

| Score             | Generator $G(p) = S(p; p)$ | Bregman divergence $D_G(q, p)$    |
|-------------------|----------------------------|-----------------------------------|
| logarithmic       | $\sum_i p_i \log p_i$      | $\text{KL}(q \  p)$               |
| Brier (quadratic) | $\ p\ _2^2$                | $\ q - p\ _2^2$                   |
| spherical         | $\ p\ _2$                  | $\ q\ _2 (1 - \cos \theta_{p,q})$ |

For the spherical row:  $\nabla G(p) = p/\|p\|$ , so  $D_G(q, p) = \|q\| - \langle p, q \rangle / \|p\| = \|q\|(1 - \cos \theta_{p,q})$ , an angular term scaled by  $\|q\|$ . For the logarithmic row the representation gives  $S(p, i) = \log p_i$  on the relative interior, with  $-\infty$  on the boundary.

**Theorem 2 (scoring rule to market maker; Hanson (2007), Abernethy et al. (2013)).** *Let  $R$  be closed, proper, and strictly convex on  $\Delta$  (in the sequel,  $R = G$ , the generator of Theorem 1 read as a regulariser), and define*

$$C(\mathbf{q}) = \sup_{p \in \Delta} (\langle p, \mathbf{q} \rangle - R(p)).$$

*Then:*

(i)  $C$  is convex and finite, the maximiser  $p^*(\mathbf{q})$  is unique, and  $\nabla C(\mathbf{q}) = p^*(\mathbf{q}) \in \Delta$ : prices are a probability vector. (Without strict convexity, prices live in  $\partial C(\mathbf{q})$ .)

(ii)  $C$  is translation-equivariant,  $C(\mathbf{q} + \alpha \mathbf{1}) = C(\mathbf{q}) + \alpha$ ; costs telescope over any trade path, so round trips cost zero and buying the full bundle costs its payout: the market is arbitrage-free.

(iii) With initial state  $\mathbf{q}_0 = \mathbf{0}$ , the maker's loss when the market settles on outcome  $i$  after net sales  $\mathbf{q}$  is

$$\text{loss}_i(\mathbf{q}) = q_i - C(\mathbf{q}) + C(\mathbf{0}) \leq R(e_i) - \inf_{p \in \Delta} R(p) \leq \sup_{p \in \Delta} R(p) - \inf_{p \in \Delta} R(p).$$

(iv) Taking  $R(p) = b \sum_i p_i \log p_i$  gives  $C(\mathbf{q}) = b \log \sum_i e^{q_i/b}$ , Hanson's LMSR, with worst-case loss  $b \log n$ .

**Proof.** (i)  $C$  is a supremum of affine functions of  $\mathbf{q}$ , hence convex; the supremum of a continuous function over the compact  $\Delta$  is attained, and strict concavity of  $p \mapsto \langle p, \mathbf{q} \rangle - R(p)$  makes the maximiser unique. Danskin's theorem gives  $\nabla C(\mathbf{q}) = p^*(\mathbf{q})$ . (ii)  $\langle p, \mathbf{q} + \alpha \mathbf{1} \rangle = \langle p, \mathbf{q} \rangle + \alpha$  for  $p \in \Delta$ , so the supremum shifts by  $\alpha$ . A trader moving the state  $\mathbf{q} \rightarrow \mathbf{q}'$  pays  $C(\mathbf{q}') - C(\mathbf{q})$  by definition, so costs over any path telescope; a round trip costs zero, and by translation equivariance the bundle  $\alpha \mathbf{1}$  costs exactly  $\alpha$ , its sure payout. (iii) The maker collects  $C(\mathbf{q}) - C(\mathbf{0})$  and pays  $q_i$ , so  $\text{loss}_i = q_i - C(\mathbf{q}) + C(\mathbf{0})$ . By Fenchel-Moreau,  $\sup_{\mathbf{q}} (\langle e_i, \mathbf{q} \rangle - C(\mathbf{q})) = R^{**}(e_i) = R(e_i)$ , and  $C(\mathbf{0}) = \sup_p -R(p) = -\inf_p R(p)$ . The final inequality holds because  $e_i \in \Delta$ ; for convex  $R$  the supremum over  $\Delta$  is attained at an extreme point, so it is typically an equality. (iv) Lagrange: maximising  $\langle p, \mathbf{q} \rangle - b \sum p_i \log p_i$  subject to  $\sum p_i = 1$  gives  $q_i - b(\log p_i + 1) = \lambda$ , so  $p_i \propto e^{q_i/b}$ , and substituting back yields  $C(\mathbf{q}) = b \log \sum_i e^{q_i/b}$ . Then  $R(e_i) = 0$  and  $\inf R = -b \log n$  at the uniform distribution, so the loss bound is  $b \log n$ . ■

**Proposition 3 (cost-function markets and CFMMs).** *Under the monotonicity, concavity, and reserve-domain hypotheses of Frongillo et al. (2024), cost-function prediction markets and constant-function market makers can be converted into one another. In the unconstrained-reserve sign convention a cost function  $C$  gives a concave CFMM potential by  $\varphi(\mathbf{r}) = -C(-\mathbf{r})$ ; bounded-reserve versions require a perspective (level-set) construction. The equivalence is a convex level-set duality; it is not the bare Fenchel conjugacy  $C \mapsto C^*$ .*

We do not reprove the general equivalence here; Frongillo et al. (2024) give it in full, and Angeris et al. (2023) and Angeris et al. (2021) develop the CFMM side.

**Example (constant product).** Take two outcome tokens with reserves  $r_1, r_2$  and invariant  $r_1 r_2 = k$ . The pool's portfolio value at prices  $(p, 1 - p)$  is

$$V(p) = \inf\{p r_1 + (1 - p) r_2 : r_1 r_2 = k\} = 2\sqrt{k p(1 - p)},$$

by the AM-GM inequality, with the infimum attained at  $r_1 = \sqrt{k(1-p)/p}$ .  $V$  is concave; reading  $R(p) = -V(p)$  as the regulariser of Theorem 2 gives a maker whose worst-case loss is the range of  $R$  on  $[0, 1]$ , namely  $\sqrt{k}$  (between  $p = \frac{1}{2}$  and the endpoints): the geometric mean of the initial reserves. The corresponding generator is not entropic, which is the convex-analytic content of the observation that constant-product markets and the LMSR price bounded-payout claims differently.

**Proposition 4 (the probability integral transform; Rosenblatt (1952), Dawid (1984)).** *If  $X$  has continuous CDF  $F$  then  $U = F(X) \sim \text{Uniform}(0, 1)$ , and  $z = \Phi^{-1}(U) \sim N(0, 1)$ . In the prequential setting, if  $F_t$  is the true conditional law of  $X_t$  given the past, then  $U_t = F_t(X_t)$  is conditionally uniform and the PIT stream is iid uniform.*

**Proof.** For continuous  $F$ , with the generalized inverse  $F^-(u) = \inf\{x : F(x) \geq u\}$ , the event  $\{F(X) \leq u\}$  differs from  $\{X \leq F^-(u)\}$  by a set of probability zero, and  $\Pr(X \leq F^-(u)) = F(F^-(u)) = u$  by continuity. The prequential statement applies this conditionally at each  $t$ . ■

Two scope remarks. First, the converse fails in the strong sense: marginally uniform PITs do not certify an informative forecast. Reporting the unconditional law  $F$  of an iid sequence gives  $F(X_t) \sim \text{Uniform}(0, 1)$  even if valuable covariates were ignored. A PIT critic therefore witnesses miscalibration relative to its test class; it complements, and does not replace, a proper score that rewards sharp conditional distributions. Conformal prediction lives on exactly this distinction: a split-conformal predictor achieves marginal coverage by construction, the uniform-PIT guarantee, while free to ignore the conditional information in the input. For a single-shape conformal predictor the wasted log score equals the mutual information  $I(R; X)$  between residual and input, and an entrant to the parimutuel pool of the companion paper who conditions on  $X$  collects it as bankroll growth at exactly that rate; the flat-in- $X$  conformal report is the break-even crowd (Theorem 9). Second, for discrete forecasts the randomized PIT preserves the exact uniform null, while the mid-PIT is a convenient deterministic diagnostic with a different null distribution.

## 4 Operators on mechanisms

**The common signature.** A *transducer* (Mealy machine) is a tuple  $(S, A, B, \delta, \lambda)$  with state space  $S$ , input alphabet  $A$ , output alphabet  $B$ , transition map  $\delta : S \times A \rightarrow S$ , and output map  $\lambda : S \times A \rightarrow B$ ; run on an input stream it produces  $s_{t+1} = \delta(s_t, a_t)$  and  $b_t = \lambda(s_t, a_t)$ , a causal map of input streams to output streams. The mechanisms of §3 share one instantiation. Let  $\text{Dist}$  denote the set of distributional beliefs and take  $S$  the wealth states,  $A = \text{Dist}^m \times \mathcal{X}$  ( $m$  participant reports and, where the stage settles, a realized outcome), and  $B = \text{Dist} \times \mathbb{R}^m$  (an aggregate belief and transfers). A *stage* is such a transducer, written

$$M : (\text{Dist}^m, w, x) \mapsto (\text{Dist}, w', \pi).$$

A scoring rule is the transfer component  $\lambda$  of a stage, not itself a map  $\text{Dist} \rightarrow \text{Dist}$ ; a market maker is a stage whose state is the inventory vector and whose emitted belief is the price; an opinion pool is a stage with no outcome argument and zero transfers. Composition wires the belief output of one stage to the belief inputs of the next while state and transfers thread through.

The operators below act on stages.

**Sequentialise.** Theorem 2: a proper score run sequentially against a wealth state is a cost-function market maker.

**Pool.** A proper score gives a batch elicitation mechanism when reports are scored independently and funded externally: properness is inherited report by report. Parimutuel and budget-balanced versions are a different game, because the pot split couples payoffs through the denominator; the companion paper’s §2 gives the price-taking analysis (truthful all-in, a symmetric equilibrium at fractional stakes, degenerate as the stake fraction vanishes), and beyond that the equilibrium theory is open.

**Ensemble (Proposition 5: the two pools are the two KL barycenters).** Let  $q_1, \dots, q_m$  be densities and  $w_i \geq 0$ ,  $\sum w_i = 1$ . Then the linear pool minimizes the forward divergences and the logarithmic pool the reverse:

$$\arg \min_p \sum_i w_i \text{KL}(q_i \| p) = \sum_i w_i q_i, \quad \arg \min_p \sum_i w_i \text{KL}(p \| q_i) \propto \prod_i q_i^{w_i}.$$

**Proof.** For the first,  $\sum_i w_i \text{KL}(q_i \| p) = \text{const} - \int (\sum_i w_i q_i) \log p$ , and  $\int \bar{q} \log p$  is maximized over densities at  $p = \bar{q}$  (Gibbs). For the second,  $\sum_i w_i \text{KL}(p \| q_i) = \int p \log p - \int p \overline{\log q}$  with  $\overline{\log q} = \sum_i w_i \log q_i$ ; the Lagrange condition is  $\log p = \overline{\log q} + \text{const}$ . ■

Training weights by cumulative log score and then mixing linearly,  $\sum_m w_m F_m$ , is Bayesian model averaging, a linear pool with score-trained weights; the logarithmic pool multiplies densities and renormalizes. They are different aggregates with different sharpness (Genest and Zidek 1986).

**Merge (Proposition 6: merging makers is infimal convolution; Rockafellar (1970), Bhaskara et al. (2023)).** For closed proper convex  $f, g$ :  $(f \square g)^* = f^* + g^*$ . Consequently, merging two cost-function makers with regularisers  $R_1, R_2$  (cost functions  $C_i = R_i^*$ ) yields the maker with regulariser  $R_1 + R_2$ , and merging  $\text{LMSR}_{b_1}$  with  $\text{LMSR}_{b_2}$  yields  $\text{LMSR}_{b_1+b_2}$ : liquidity adds.

**Proof.**  $(f \square g)^*(p) = \sup_x \langle p, x \rangle - \inf_y \{f(y) + g(x - y)\} = \sup_{y,z} \langle p, y \rangle - f(y) + \langle p, z \rangle - g(z) = f^*(p) + g^*(p)$ . For the makers,  $C_1 \square C_2 = (R_1 + R_2)^*$  by biconjugacy, and  $b_1 G + b_2 G = (b_1 + b_2)G$  for the entropic generator. ■

This is also the risk-sharing literature’s composition law: the aggregate of several agents’ convex risk measures is their infimal convolution, with the Pareto allocation as minimiser and the common subgradient as the clearing price (Barrieu and El Karoui 2005; Jouini et al. 2008); the market-liquidity reading is developed by Bhaskara et al. (2023).

**Conjugate.** Run a stage in a transformed coordinate and map back. This is where propriety needs care, and the theory splits in two:

- *Fixed transformations.* Properness is preserved under any fixed transformation of forecast and outcome, and strict propriety iff the transformation is injective (Allen et al. 2023, Prop. 4; Pic et al. 2025, Prop. 1). The Markov-kernel (stochastic channel) extension, strict propriety iff the channel is injective on laws, is Theorem 2 of the companion point-cloud paper, whose Theorem 1 also exhibits the canonical failure: a KDE smoothing seam whose raw-outcome score elicits a deconvolution.
- *Forecast-dependent transformations.* The PIT critic transforms outcomes by the reported  $F$  itself, so the fixed-transformation theorems do not apply to it; Proposition 4 and its scope remarks are the correct warrant, and the critic is a calibration diagnostic rather than a strictly proper score. The stacked-lottery design (Cotton 2020, slides 29-31) is this operator in practice: percentiles from one game feed the next, and calibration is produced by composing monotone maps contributed by competing algorithms.

**Residual.** Let a stage emit the aggregate  $F_1$ , and let a second market elicit a distribution for the residual  $U = F_1(X)$ , settling at the realised  $u = F_1(x)$ . If  $F_1$  is the true conditional law then  $U$  is uniform (Proposition 4) and the second market has nothing to price; whatever structure remains in the residual is the second stage’s edge. The corrected forecast composes the two reports.

**Proposition 7 (the correction is a multiplicative reweighting).** *Let  $F_1$  be strictly increasing with density  $p_1 > 0$  and let the residual market’s consensus be a distribution  $G$  on  $[0, 1]$  with density  $g$ . The composed forecast  $F = G \circ F_1$  has density*

$$p(x) = p_1(x) g(F_1(x)),$$

*so  $\log p(x) = \log p_1(x) + \log g(u)$  with  $u = F_1(x)$ : the chain’s log score is the sum of stage log scores, and the residual stage is paid by a proper score on  $u$  alone.*

**Proof.** Chain rule:  $F'(x) = g(F_1(x)) p_1(x)$ ; take logarithms. The residual score  $\log g(u)$  is the logarithmic score of Theorem 1 applied to the report  $g$  and outcome  $u$ , hence strictly proper for the law of  $U$ . ■

Multiplying the density by a ratio fitted to what the current model gets wrong is the functional-gradient step of boosting under log loss (Mason et al. 1999; Friedman 2001), so a chain of residual markets is stagewise boosting with wealth as the learning rate. What Proposition 7 does not settle is the game across stages: who funds the residual pot, and whether a forecaster free to enter both stages prefers to withhold information from the first and sell it to the second (§9).

**Spec.** Serialise a pipeline to data and search over it; the mechanism analogue is a market over pipelines. Also open.

**Stagewise play.** Call a profile of reports a *stagewise equilibrium* if no participant gains by a deviation confined to a single stage, all other stages’ reports held fixed. This is the pipeline version of the myopic-trader assumption standard in the market-scoring-rule literature (Hanson 2007; Chen and Vaughan 2010).

**Proposition 8 (single-stage guarantees compose under stagewise play).** *Suppose each stage of a pipeline, taken in isolation with its inputs fixed, makes the truthful report a best response: strict propriety for externally funded scoring stages (Theorem 1), the sequential scoring of Theorem 2 for market stages, the price-taking pot-split analysis of the companion paper for parimutuel stages, and the residual score of Proposition 7 for correction stages. Then truthful reporting at every stage is a stagewise equilibrium of the pipeline.*

**Proof.** A deviation confined to stage  $k$  changes the deviator’s payoff only through stage  $k$ ’s transfer map, since every other stage’s reports, and hence its inputs, are held fixed; the stage- $k$  hypothesis then makes the truthful report a best response. ■

What the concept excludes is cross-stage deviation: a downstream stake is a derivative written on an upstream settlement, with the attendant incentives to distort the underlying (Kumar and Seppi 1992; Jarrow 1994; Hanson and Oprea 2009; Ostrovsky 2012). The note works under stagewise play throughout.

## 5 Betting against a conformal predictor

The introduction claimed that split-conformal prediction is a degenerate composition and that running the residual stage as a market collects what the degeneracy wastes. This section proves it. A point predictor leaves a residual  $R$ ; conformalization re-levels its marginal law. Write

$U = G(R)$  for the PIT of the residual,  $G$  the marginal CDF of  $R$ , so  $U$  is marginally uniform whatever the model (Proposition 4), and let  $g(u \mid x)$  be the conditional density of  $U$  given  $X = x$ . Since  $R \mapsto G(R)$  is almost surely invertible,  $I(U; X) = I(R; X)$ , and because the marginal of  $U$  is uniform,

$$I(R; X) = \mathbb{E}_X \text{KL}(P_{R|X} \| G) = \mathbb{E}_X \int_0^1 g(u \mid X) \log g(u \mid X) du.$$

The residual stage is the parimutuel of the companion paper run on the rank scale. Two facts about that pool are needed.

**Lemma 1 (pool payoff).** *Normalise the pool to one and let the crowd's aggregate stake have density  $q$  on the outcome space. An infinitesimal, price-taking entrant who stakes unit wealth with density  $b$  holds, after outcome  $u$ , the wealth  $W(u) = b(u)/q(u)$ .*

**Proof.** Bin the outcome axis at width  $\delta$ . The crowd stakes  $q(u)\delta$  of the unit pool on the bin at  $u$ , the entrant  $b(u)\delta$ . If the outcome lands there the pool is split in proportion to stake, so the per-unit payoff is  $1/(q(u)\delta)$  and the entrant collects  $b(u)\delta \cdot 1/(q(u)\delta) = b(u)/q(u)$ . The bin width cancels and the limit is the Radon-Nikodym derivative  $db/dq$ . ■

The cancellation is what makes a lottery on a point outcome well posed: the pool divides two vanishing quantities and the ratio survives, with no reference measure and no posted odds.

**Lemma 2 (log-optimal growth).** *If the outcome is drawn from  $p$ , the expected log-growth of an entrant staking  $b$  against a crowd staking  $q$  is*

$$\mathbb{E}_{U \sim p} \left[ \log \frac{b(U)}{q(U)} \right] = \text{KL}(p \| q) - \text{KL}(p \| b) \leq \text{KL}(p \| q),$$

*with equality iff  $b = p$ : the log-optimal stake is the truth.*

**Proof.** Add and subtract  $\log p$  inside the expectation; Gibbs' inequality kills the second term exactly at  $b = p$ . ■

The lemmas are the parimutuel form of the Kelly-Breiman log-optimality principle (Kelly 1956; Cover and Thomas 2006; Barron and Cover 1988); when the crowd prices the marginal, the growth of belief  $b$  is  $I(R; X) - \mathbb{E}_X \text{KL}(g(\cdot \mid X) \| b(\cdot \mid X))$ , a decomposition written by Kemp and Bettencourt (2022) for population growth in stochastic environments.

**Theorem 9 (parimutuel rent of a conformal predictor).** *Put the pool on the rank scale. The single-shape conformal predictor is the crowd that prices it flat,  $q \equiv 1$  on  $[0, 1]$ . Against it:*

(i) *an entrant who knows only the marginal stakes  $b \equiv 1$  and grows at rate zero: marginal coverage stated as wealth;*

(ii) *an entrant who observes  $X$  and stakes  $b = g(\cdot \mid X)$  grows per round at rate  $\mathbb{E}_X \text{KL}(g(\cdot \mid X) \| 1) = I(R; X)$ ;*

(iii) *with a track take  $\tau \in [0, 1]$  the informed rate is  $I(R; X) + \log(1 - \tau)$ , positive iff  $I(R; X) > -\log(1 - \tau)$ .*

**Proof.** Lemma 1 with  $q \equiv 1$  makes the payoff  $b(U)$ ; Lemma 2 makes the growth  $\text{KL}(p \| 1) - \text{KL}(p \| b)$  for true conditional law  $p = g(\cdot \mid x)$ . For (i),  $b = 1$  and the marginal of  $U$  is uniform, so the rate is zero. For (ii), conditioning on  $X = x$  the optimal stake is  $g(\cdot \mid x)$  with conditional rate  $\text{KL}(g(\cdot \mid x) \| 1) = \int g \log g$ ; averaging over  $X$  gives  $I(R; X)$  by the display above. For (iii) every payoff is multiplied by  $1 - \tau$ , adding  $\log(1 - \tau)$  to the rate. ■

The certificate and the leak are two readings of one pool. Marginal coverage is part (i), a break-even statement; part (ii) is what the certificate cannot see, and re-leveling cannot close

it because re-leveling acts on the marginal, where the leak is invisible. Conformal re-leveling takes no rake, so any conditional information at all is pure profit.

**Example (Gaussian residual).** Let  $(R, X)$  be jointly Gaussian with correlation  $\rho$ , marginally  $R \sim N(0, \sigma_R^2)$ ,  $X \sim N(0, \sigma_X^2)$ ; any misspecified linear predictor produces this. The informed entrant's report is  $R \mid X = x \sim N(\rho(\sigma_R/\sigma_X)x, (1 - \rho^2)\sigma_R^2)$ , and with  $m(x) = \rho(\sigma_R/\sigma_X)x$ ,  $s^2 = (1 - \rho^2)\sigma_R^2$ ,

$$\mathbb{E}_X \text{KL}(N(m, s^2) \parallel N(0, \sigma_R^2)) = \mathbb{E}_X \left[ \log \frac{\sigma_R}{s} + \frac{s^2 + m^2}{2\sigma_R^2} - \frac{1}{2} \right] = -\frac{1}{2} \log(1 - \rho^2) = I(R; X).$$

At  $\rho = \frac{1}{2}$  the rent is 0.144 nats per observation and the informed bankroll doubles every five observations; at  $\rho = 0.3$ , every fifteen. Throughout, the conformal band's marginal coverage remains exact. The composed forecast of Proposition 7 built from the entrant's rank report  $g(u \mid x) = p_{R|X}(G^-(u))/p_R(G^-(u))$  recovers the conditional law exactly:  $p_R(r) g(G(r) \mid x) = p_{R|X}(r)$ .

**The mechanism has run in production.** The microprediction platform's nearest-the-pin pool (Cotton 2022) paid each entry its sample density at the realised value relative to the field's, the payoff of Lemma 1 with densities estimated from submitted samples; the MidOne contest (CrunchDAO 2024) priced an explicit density (Cotton 2024a) the same way, and its priced object was a residual stream. A conformal predictor entering such a pool is the participant that prices flat in  $X$ , and Theorem 9 is the bankroll of its better-informed competitor.

**Measuring the rent.** The rent is defined through a log score, so it is estimated rather than read off. Two routes:

**Proposition 10 (certified lower bound).** Let  $\text{HSIC}(U, X)$  be computed with a bounded product kernel,  $k \leq K$ . Then  $I(R; X) \geq \text{HSIC}(U, X)/(2K)$ .

**Proof.**  $\text{MMD} = \text{HSIC}^{1/2}$  is the integral probability metric over the unit ball of the tensor RKHS, whose witnesses satisfy  $\|f\|_\infty \leq \sqrt{K}$ , so  $\text{MMD} \leq 2\sqrt{K}$  TV between the joint law and the product of marginals. Pinsker's inequality (Cover and Thomas 2006) gives  $I(U; X) = \text{KL} \geq 2 \text{TV}^2$ , and  $I(U; X) = I(R; X)$ . Combining,  $I(R; X) \geq 2(\text{MMD}/2\sqrt{K})^2 = \text{HSIC}/(2K)$ . ■

Distance covariance (Székely et al. 2007) is a fixed multiple of an HSIC (Sejdinovic et al. 2013), so a measured dependence between rank and input certifies a minimum extractable rent. No matching upper bound exists with a fixed kernel: MMD metrizes weak convergence (Simon-Gabriel et al. 2023) and KL does not.

**Proposition 11 (anytime-valid test and consistent estimate).** At round  $t$  choose a betting density  $b_t(\cdot \mid x) \geq 0$  with  $\int_0^1 b_t(u \mid x) du = 1$  from the history, and set  $W_t = \prod_{s \leq t} b_s(U_s \mid X_s)$ . (i) Under  $H_0 : U \perp X$ ,  $(W_t)$  is a non-negative martingale with mean one, so  $\Pr(\sup_t W_t \geq 1/\alpha) \leq \alpha$  (Ville 1939): rejecting when  $W_t \geq 1/\alpha$  is an anytime-valid level- $\alpha$  test of conditional miscoverage. (ii) If  $b_t \rightarrow g$ , then  $t^{-1} \log W_t \rightarrow I(R; X)$  almost surely.

**Proof.** (i) Under  $H_0$ ,  $U_t \mid X_t$  is uniform and independent of the past, so each factor has conditional mean  $\int_0^1 b_t = 1$ ; Ville's inequality applies to the resulting martingale. (ii)  $t^{-1} \sum_s \log b_s(U_s \mid X_s) \rightarrow \mathbb{E}[\log g(U \mid X)] = I(R; X)$  by the law of large numbers and the display opening this section. ■

The wealth process is simultaneously the test and the estimator; it is the log-optimal e-variable of safe testing (Grünwald et al. 2024; Vovk and Wang 2021) for the null  $U \perp X$ , with

conformal test martingales as ancestor (Vovk et al. 2003). Sequential tests of forecast calibration (Arnold et al. 2023) and conditional independence by betting (Shaer et al. 2023) use the same mechanics; the identification of the growth rate with the conformal predictor’s information gap, and with the realised bankroll in a deployed pool, is the reading here. One measurement caveat:  $U = G(R)$  uses the true marginal, and with empirical ranks the null is exact in the full-conformal, leave-one-out sense. And one scope remark: by the distribution-free no-go results (Lei and Wasserman 2014; Foygel Barber et al. 2021; Vovk 2012) no procedure certifies conditional coverage from data alone, so a quiet e-process means no rent found at this power, never conditional validity.

## 6 Samples as messages

The message type has so far been a distribution given exactly. Deployed contests collect something rougher: a finite cloud of samples, smoothed by the operator into a density before settlement. The companion paper gives the sample-based elicitation result the algebra needs (Cotton 2026a):

**Proposition 12 (sample-based elicitation; Cotton (2026a), Thms 1-2).** *Score the bandwidth- $h$  kernel density estimate of a submitted cloud by the logarithmic score. Settled at the raw outcome, the optimal cloud is drawn from a deconvolution of the belief (for Gaussian beliefs and kernels, the belief with  $h^2$  removed from the variance). Settled at the outcome jittered by the same kernel, truthful sampling is optimal, and strictly so whenever the kernel’s characteristic function has dense nonzero support: the smoothing channel is injective on laws.*

The proof is in the companion paper. What matters here is what the repair buys the algebra: with jittered settlement the cloud is a valid message, and the operators of §4 act on clouds directly.

- *Conjugate* acts pointwise: the pushforward of a cloud under  $\psi$  is  $\psi$  applied to each sample, with no Jacobian computed anywhere. This is the practical reason to prefer sample messages: transformation of densities requires calculus, transformation of clouds requires arithmetic.
- *Pool* (linear) is a wealth-weighted union: sample from participant  $i$ ’s cloud with probability proportional to  $w_i$ . The linear pool is the sample-native aggregate; the logarithmic pool has no comparably clean sample form, since multiplying densities requires importance weights.
- *Residual* acts by ranking: pass each sample through the upstream CDF. Composition of monotone maps, the stacked-lottery operation, is again pointwise on samples.
- *Sequentialise* is the nearest-the-pin pool itself: the market’s state is the field’s aggregate cloud.

One caveat governs the order of operations. Smoothing and conjugation commute only for affine maps: for  $\psi(x) = ax + b$  and a Gaussian kernel, the KDE of the mapped cloud at bandwidth  $|a|h$  equals the pushforward of the bandwidth- $h$  KDE, and for nonlinear  $\psi$  no bandwidth makes this an identity. The smoothing seam and its jitter must therefore be applied in the settlement coordinates, after all transforms, which in pipeline terms says the KDE stage belongs immediately before settlement and nowhere else.

## 7 Margins and copulas

A multivariate settlement can be factored. By Sklar’s theorem (Sklar 1959) a joint law with continuous margins  $F_1, \dots, F_d$  and copula  $C$  has density

$$p(x) = \prod_{i=1}^d f_i(x_i) \cdot c(F_1(x_1), \dots, F_d(x_d)),$$

with  $c$  the copula density on  $[0, 1]^d$ .

**Proposition 13 (the log score factors).** *For a joint report assembled from margin reports  $f_i$  and a copula report  $c$ ,*

$$\log p(x) = \sum_{i=1}^d \log f_i(x_i) + \log c(u), \quad u_i = F_i(x_i),$$

*so a pipeline of  $d$  margin stages and one copula stage settled on the rank vector  $u$  pays every stage a logarithmic score of its own object, and the chain’s log score is the sum. Given the margin reports, the copula stage’s score is strictly proper for the law of the rank vector; that law is the copula of the joint iff the margins are correct.*

**Proof.** Take logarithms in Sklar’s density factorization; the summands are exactly the stage scores. Strict propriety of the rank stage is Theorem 1 applied to densities on  $[0, 1]^d$ , the transform  $x \mapsto u$  being fixed once the margin reports are given. The final clause is Sklar’s theorem applied to the true joint law. ■

This is the two-stage estimation logic of inference functions for margins (Joe 1997), and the factorization underlies out-of-sample copula comparison (Diks et al. 2010); here it is read as a market design. The copula stage is the multivariate Residual operator: under correct margins the rank vector has uniform margins and its joint law is the dependence structure alone, so a rank-settled market elicits dependence separately from level, and is invariant to any recalibration of the margins. The microprediction platform ran exactly this factoring: alongside univariate z-streams it operated bivariate and trivariate streams in which community-implied percentiles were embedded by a space-filling curve into one settled scalar (Cotton 2022), a copula market in production. Sample messages compose with the factoring: rank each coordinate of a submitted cloud through the margin CDFs and the result is a cloud on  $[0, 1]^d$  for the copula stage, with the §6 seam discipline applied at that stage’s settlement.

## 8 Implementation notes

The mechanisms of the catalogue are implemented against a single interface patterned on the skaters time-series contract (Cotton 2026b), a successor to timemachines (Cotton 2021): every stage consumes and emits the same distributional type and threads its own state, so the signature of §4 is the forecasting contract with wealth in place of model state. The closure of the message type is what makes a small operator set sufficient; the transforms, ensembles, and residual constructions of §4 all consume and emit the one type. Scores are implemented in loss form; the text uses reward form.

## 9 Open problems

1. *Conservation of edge.* Wealth conservation is automatic, a sum of stage-level zeros, and is not the invariant. The invariant that can fail is the edge a truthful participant holds over the price,  $D_G(q, \pi)$ . For the logarithmic score the edge is  $\text{KL}(q \parallel \pi)$  and the question is closed by the data-processing inequality: an interface channel cannot increase it, and preserves it exactly when the channel is sufficient for the pair (Csiszár 1967). Open is the same statement for the Bregman edge of a general generator  $G$ , and the rate at which edge leaks through an approximately non-sufficient interface.
2. *Residual markets.* Proposition 7 gives the single-step correction; a chain of residual markets is stagewise boosting with wealth as the learning rate. Open: the microstructure (who funds each residual pot, and when it settles relative to the stage before), whether a forecaster free to enter several stages prefers withholding information upstream to sell it downstream, and whether the iterated correction converges to the true conditional law as boosting does. Split-conformal prediction is the one-participant degenerate case, with its rent  $I(R; X)$  as the value left on the table (Theorem 9). The two-settlement electricity market is the live analogue in point prices, with convergence bidders as the residual stage’s informed entrants (Jha and Wolak 2023); the distributional version is the gap.
3. *Copula markets.* A rank-settled stage elicits dependence separately from margins (Proposition 13) and is invariant to margin recalibration. The racetrack’s win and exotic pools are the live analogue on a finite outcome space, margins and joint orders in separate simultaneous pools, with the Harville map as the consistency condition between them (Harville 1973; Hausch et al. 1981). Open: the equilibrium when the same participants trade margin and copula stages, and the choice of embedding for  $d \geq 2$ , the space-filling curves as deployed against the random one-dimensional projections of the companion paper.

## References

- Abernethy, Jacob, Yiling Chen, and Jennifer Wortman Vaughan. 2013. “Efficient Market Making via Convex Optimization, and a Connection to Online Learning.” *ACM Transactions on Economics and Computation* 1 (2): 12:1–39. <https://doi.org/10.1145/2465769.2465777>.
- Allen, Sam, David Ginsbourger, and Johanna Ziegel. 2023. “Evaluating Forecasts for High-Impact Events Using Transformed Kernel Scores.” *SIAM/ASA Journal on Uncertainty Quantification* 11 (3): 906–40. <https://doi.org/10.1137/22M1532184>.
- Angeris, Guillermo, Tarun Chitra, Alex Evans, and Matthew Lorig. 2023. “Short Communication: A Primer on Perpetuals.” *SIAM Journal on Financial Mathematics* 14 (1): SC17–30. <https://doi.org/10.1137/22M1520931>.
- Angeris, Guillermo, Hsien-Tang Kao, Rei Chiang, Charlie Noyes, and Tarun Chitra. 2021. “An Analysis of Uniswap Markets.” *Cryptoeconomic Systems* 1 (1). <https://arxiv.org/abs/1911.03380>.
- Arnold, Sebastian, Alexander Henzi, and Johanna F. Ziegel. 2023. “Sequentially Valid Tests for Forecast Calibration.” *The Annals of Applied Statistics* 17 (3): 1909–35.

- Atanasov, Pavel, Phillip Rescober, Eric Stone, et al. 2017. "Distilling the Wisdom of Crowds: Prediction Markets Vs. Prediction Polls." *Management Science* 63 (3): 691–706.
- Banerjee, Arindam, Xin Guo, and Hui Wang. 2005. "On the Optimality of Conditional Expectation as a Bregman Predictor." *IEEE Transactions on Information Theory* 51 (7): 2664–69. <https://doi.org/10.1109/TIT.2005.850145>.
- Barrieu, Pauline, and Nicole El Karoui. 2005. "Inf-Convolution of Risk Measures and Optimal Risk Transfer." *Finance and Stochastics* 9 (2): 269–98. <https://doi.org/10.1007/s00780-005-0152-0>.
- Barron, Andrew R., and Thomas M. Cover. 1988. "A Bound on the Financial Value of Information." *IEEE Transactions on Information Theory* 34 (5): 1097–100.
- Bhaskara, Adithya, Rafael Frongillo, and Maneesha Papireddygar. 2023. *A General Theory of Liquidity Provisioning for Prediction Markets*. arXiv:2311.08725. <https://arxiv.org/abs/2311.08725>.
- Chen, Yiling, and Jennifer Wortman Vaughan. 2010. "A New Understanding of Prediction Markets via No-Regret Learning." *Proceedings of the 11th ACM Conference on Electronic Commerce (EC)*. <https://doi.org/10.1145/1807342.1807372>.
- Cotton, Peter. 2020. *The Lottery Paradox: A New Use*. <https://www.slideshare.net/slideshow/1-lottery-paradox-csaildec2020/242173597>.
- Cotton, Peter. 2021. *Timemachines: Time-Series Prediction with a Uniform Skater Interface*. <https://github.com/microprediction/timemachines>.
- Cotton, Peter. 2022. *Microprediction: Building an Open AI Network*. MIT Press.
- Cotton, Peter. 2024a. *Density: Conventions for Densities Used in Various Games*. <https://github.com/microprediction/density>.
- Cotton, Peter. 2024b. *Monteprediction: A Weekly Monte Carlo Competition in Eleven Dimensions*. <https://www.monteprediction.com>.
- Cotton, Peter. 2026a. *Scoring Point-Cloud Distributional Submissions*. <https://mechanisms.microprediction.org/papers/scoring-point-cloud-distributional-submissions.html>.
- Cotton, Peter. 2026b. *Skaters: Fast Online Time-Series Models with Composable Transforms*. <https://github.com/microprediction/skaters>.
- Cover, Thomas M., and Joy A. Thomas. 2006. *Elements of Information Theory*. 2nd ed. Wiley.
- Craib, Richard, Geoffrey Bradway, Xander Dunn, and Joey Krug. 2017. *Numeraire: A Cryptographic Token for Coordinating Machine Intelligence*. <https://numer.ai/whitepaper.pdf>.

- Cramton, Peter. 2017. "Electricity Market Design." *Oxford Review of Economic Policy* 33 (4): 589–612.
- CrunchDAO. 2024. *The MidOne Contest*. <https://mid-one.crunchdao.com/>.
- Csiszár, Imre. 1967. "Information-Type Measures of Difference of Probability Distributions and Indirect Observations." *Studia Scientiarum Mathematicarum Hungarica* 2: 299–318.
- Dawid, A. P. 1984. "Present Position and Potential Developments: Some Personal Views. Statistical Theory: The Prequential Approach." *Journal of the Royal Statistical Society. Series A (General)* 147 (2): 278–92. <https://doi.org/10.2307/2981683>.
- Diks, Cees, Valentyn Panchenko, and Dick van Dijk. 2010. "Out-of-Sample Comparison of Copula Specifications in Multivariate Density Forecasts." *Journal of Economic Dynamics and Control* 34 (9): 1596–609.
- Foygel Barber, Rina, Emmanuel J. Candès, Aaditya Ramdas, and Ryan J. Tibshirani. 2021. "The Limits of Distribution-Free Conditional Predictive Inference." *Information and Inference* 10 (2): 455–82.
- Friedman, Jerome H. 2001. "Greedy Function Approximation: A Gradient Boosting Machine." *The Annals of Statistics* 29 (5): 1189–232. <https://doi.org/10.1214/aos/1013203451>.
- Frongillo, Rafael, Maneesha Papireddygar, and Bo Waggoner. 2024. "An Axiomatic Characterization of CFMMs and Equivalence to Prediction Markets." *15th Innovations in Theoretical Computer Science Conference (ITCS), LIPIcs*, vol. 287: 51:1–21. <https://doi.org/10.4230/LIPIcs.ITCS.2024.51>.
- Genest, Christian, and James V. Zidek. 1986. "Combining Probability Distributions: A Critique and an Annotated Bibliography." *Statistical Science* 1 (1): 114–35. <https://doi.org/10.1214/ss/1177013825>.
- Gneiting, Tilmann, and Adrian E. Raftery. 2007. "Strictly Proper Scoring Rules, Prediction, and Estimation." *Journal of the American Statistical Association* 102 (477): 359–78. <https://doi.org/10.1198/016214506000001437>.
- Grünwald, Peter, Rianne de Heide, and Wouter M. Koolen. 2024. "Safe Testing." *Journal of the Royal Statistical Society B* 86 (5): 1091–128.
- Hanson, Robin. 2007. "Logarithmic Market Scoring Rules for Modular Combinatorial Information Aggregation." *The Journal of Prediction Markets* 1 (1): 3–15. <https://doi.org/10.5750/jpm.v1i1.417>.
- Hanson, Robin, and Ryan Oprea. 2009. "A Manipulator Can Aid Prediction Market Accuracy." *Economica* 76 (302): 304–14. <https://doi.org/10.1111/j.1468-0335.2008.00734.x>.

- Harville, David A. 1973. "Assigning Probabilities to the Outcomes of Multi-Entry Competitions." *Journal of the American Statistical Association* 68 (342): 312–16.
- Hausch, Donald B., William T. Ziemba, and Mark Rubinstein. 1981. "Efficiency of the Market for Racetrack Betting." *Management Science* 27 (12): 1435–52.
- Jarrow, Robert A. 1994. "Derivative Security Markets, Market Manipulation, and Option Pricing Theory." *Journal of Financial and Quantitative Analysis* 29 (2): 241–61. <https://doi.org/10.2307/2331224>.
- Jha, Akshaya, and Frank A. Wolak. 2023. "Can Financial Participants Improve Price Discovery and Efficiency in Multi-Settlement Markets with Trading Costs?" *American Economic Journal: Economic Policy* 15 (2): 301–40.
- Joe, Harry. 1997. *Multivariate Models and Dependence Concepts*. Chapman & Hall.
- Jouini, Elyes, Walter Schachermayer, and Nizar Touzi. 2008. "Optimal Risk Sharing for Law Invariant Monetary Utility Functions." *Mathematical Finance* 18 (2): 269–92. <https://doi.org/10.1111/j.1467-9965.2007.00332.x>.
- Kelly, Jr., J. L. 1956. "A New Interpretation of Information Rate." *The Bell System Technical Journal* 35 (4): 917–26. <https://doi.org/10.1002/j.1538-7305.1956.tb03809.x>.
- Kemp, Jordan T., and Luís M. A. Bettencourt. 2022. *The Bayesian Origins of Growth Rates in Stochastic Environments*. arXiv:2209.09492. <https://arxiv.org/abs/2209.09492>.
- Kumar, Praveen, and Duane J. Seppi. 1992. "Futures Manipulation with 'Cash Settlement'." *The Journal of Finance* 47 (4): 1485–502. <https://doi.org/10.1111/j.1540-6261.1992.tb04665.x>.
- Lei, Jing, and Larry Wasserman. 2014. "Distribution-Free Prediction Bands for Non-Parametric Regression." *Journal of the Royal Statistical Society B* 76 (1): 71–96.
- Mason, Llew, Jonathan Baxter, Peter L. Bartlett, and Marcus Frean. 1999. "Boosting Algorithms as Gradient Descent." *Advances in Neural Information Processing Systems (NeurIPS)*, 512–18.
- McCarthy, John. 1956. "Measures of the Value of Information." *Proceedings of the National Academy of Sciences* 42 (9): 654–55. <https://doi.org/10.1073/pnas.42.9.654>.
- Ostrovsky, Michael. 2012. "Information Aggregation in Dynamic Markets with Strategic Traders." *Econometrica* 80 (6): 2595–647. <https://doi.org/10.3982/ECTA8479>.
- Pic, Romain, Clément Dombry, Philippe Naveau, and Maxime Taillardat. 2025. "Proper Scoring Rules for Multivariate Probabilistic Forecasts Based on Aggregation and Transformation." *Advances in Statistical Climatology, Meteorology and Oceanography* 11 (1): 23–58. <https://doi.org/10.5194/asmo-11-23-2025>.

- Rockafellar, R. Tyrrell. 1970. *Convex Analysis*. Princeton Mathematical Series 28. Princeton University Press. <https://doi.org/10.1515/9781400873173>.
- Rosenblatt, Murray. 1952. "Remarks on a Multivariate Transformation." *The Annals of Mathematical Statistics* 23 (3): 470–72. <https://doi.org/10.1214/aoms/1177729394>.
- Savage, Leonard J. 1971. "Elicitation of Personal Probabilities and Expectations." *Journal of the American Statistical Association* 66 (336): 783–801. <https://doi.org/10.1080/01621459.1971.10482346>.
- Sejdinovic, Dino, Bharath Sriperumbudur, Arthur Gretton, and Kenji Fukumizu. 2013. "Equivalence of Distance-Based and RKHS-Based Statistics in Hypothesis Testing." *The Annals of Statistics* 41 (5): 2263–91. <https://doi.org/10.1214/13-AOS1140>.
- Shaer, Shalev, Gal Maman, and Yaniv Romano. 2023. "Model-X Sequential Testing for Conditional Independence via Testing by Betting." *AISTATS*. <https://arxiv.org/abs/2210.00354>.
- Simon-Gabriel, Carl-Johann, Alessandro Barp, Bernhard Schölkopf, and Lester Mackey. 2023. "Metρίζing Weak Convergence with Maximum Mean Discrepancies." *Journal of Machine Learning Research* 24.
- Sklar, Abe. 1959. "Fonctions de répartition à  $n$  Dimensions Et Leurs Marges." *Publications de l'Institut de Statistique de l'Université de Paris* 8: 229–31.
- Székely, Gábor J., Maria L. Rizzo, and Nail K. Bakirov. 2007. "Measuring and Testing Dependence by Correlation of Distances." *The Annals of Statistics* 35 (6): 2769–94.
- Ville, Jean. 1939. *Étude Critique de La Notion de Collectif*. Gauthier-Villars.
- Vovk, Vladimir. 2012. "Conditional Validity of Inductive Conformal Predictors." *Asian Conference on Machine Learning*, PMLR, vol. 25: 475–90.
- Vovk, Vladimir, Ilia Nouretdinov, and Alex Gammerman. 2003. "Testing Exchangeability on-Line." *International Conference on Machine Learning*.
- Vovk, Vladimir, and Ruodu Wang. 2021. "E-Values: Calibration, Combination and Applications." *The Annals of Statistics* 49 (3): 1736–54.